

# Higher Dimensional Collaborative Filtering for Beer Networks

Chris Jerrett, Athreya Murali, Ujwal Pandey  
*Computer Science Department*  
*Rensselaer Polytechnic Institute*  
 Troy, NY, USA

## Abstract

We introduce a new method for collaborative filtering for recommender systems with vector weight edges. In this paper we detail how to compute tensor decomposition for a special adjacency tensors and detail the results.

## I. INTRODUCTION

Recommender systems have been widely used for creating personalized predictive models that help individuals identify content of interest. Collaborative filtering (CF) is the '*de-facto*' standard used for these systems. Matrix Factorization (MF) is a popular method to perform CF and has been used in many popular systems i.e. Netflix Movie Prize [1]. The general problem set-up considers a set of users  $U = \{u_1, u_2, \dots, u_n\}$ , a set of reviews of a product  $\{r_1, r_2, \dots, r_m\}$  and set of ratings for the product  $\{w_1, w_2, \dots, w_m\}$  where  $w_i \in \mathbb{R}$ . This set-up can be viewed as a bipartite graph  $G_{X,Y} = (V, E)$  and the goal is to predict edges and their weights. Using MF, we consider the adjacency matrix  $A \in \mathbb{R}^{n \times m}$  to solve the following problem

$$\min_{U,V} \|A - UV^T\|_F^2$$

where  $U \in \mathbb{R}^{n \times k}$  is the user-feature matrix and  $V \in \mathbb{R}^{m \times k}$  is the product-feature matrix. There are many methods to solve this problem such as Singular-Value Decomposition (SVD), Stochastic Gradient Descent (SGD), and Weighted Alternating Least-Squares (WALS). However, issues of sparsity of the dataset tend to make this problem a little more difficult to solve.

In this paper, we consider a variant of this problem; we consider a product with a multi-attribute rating. The problem changes slightly since our set of ratings  $\{w_1, w_2, \dots, w_m\}$  now has  $w_i \in \mathbb{R}^k$ . To model this problem, we now consider the adjacency tensor  $\mathcal{A} \in \mathbb{R}^{n \times m \times k}$ , a real 3-dimensional array, where its element is denoted by  $\mathcal{A}_{u_i, r_i} = w_i$ . In this paper, we explore mathematically and experiment with Masked CP-Decomposition, a tensor decomposition method, on the BeerAdvocate dataset provided by SNAP Labs. We discuss the pre-processing of the dataset, capture the experiments by measuring the error of method using both proximal gradient descent and stochastic gradient descent, and compare our results to a simplistic model.

## II. BACKGROUND

We start by presenting a few definitions required to understand this work [2].

**Definition II.1** (Simple Tensor). A simple tensor or order  $n$  is a tensor  $\mathcal{X}$  such that there exists  $n$  vectors  $X_1, X_2, \dots, X_n$  such that  $\mathcal{X} = X_1 \otimes X_2 \otimes \dots \otimes X_n$ . Where  $\otimes$  is the outer product.

**Definition II.2** (Tensor Rank). The rank of a tensor  $\mathcal{X}$  is the minimum number of simple tensors  $\mathcal{X}_1, \mathcal{X}_2, \dots, \mathcal{X}_F$  such that  $\mathcal{X} = \mathcal{X}_1 + \mathcal{X}_2 + \dots + \mathcal{X}_F$ .

**Definition II.3** (CP Decomposition). The CANDECOMP/PARAFAC (CP) decomposition of an order  $N$  tensor  $\mathcal{X}$  is a tensor rank  $F$  approximation of  $\mathcal{X}$  such that,

$$\begin{aligned}\mathcal{X} &\approx \llbracket U_1, U_2, \dots, U_N \rrbracket \\ &= \sum_{i=1}^F U_1(:, i) \otimes U_2(:, i) \otimes \dots \otimes U_N(:, i)\end{aligned}$$

where  $U_i(:, j)$  is the  $j$ th column of the  $i$ th feature matrix.

**Definition II.4** (Mode- $n$  unfolding). The mode- $n$  unfolding of a tensor  $\mathcal{X} \in \mathbb{R}^{I_1 \times I_2 \times \dots \times I_N}$  is a matrix  $X \in \mathbb{R}^{I_n \times I_1 I_2 \dots I_{n-1} I_{n+1} \dots I_N}$  such that

$$\mathcal{X}(i_1, i_2, \dots, i_N) = X_{(n)}(j, i_n)$$

with  $j = 1 + \sum_{k=1, k \neq n}^N (i_k - 1)j_k$  and  $j = \prod_{m=1, m \neq k}^{k-1} I_m$ .

We note that the mode  $n$  unfolding of a tensor is a way to effectively turn the tensor into a matrix by concatenating mode- $n$  slices.

**Definition II.5** (Mode- $n$  fiber). The mode- $n$  fiber is the vector obtained by fixing all indices except for the  $i$ -th index.

$$\mathcal{X}(i_1, i_2, \dots, i_{n-1}, :, i_{n+1}, \dots, i_N)$$

We can describe the dataset using a weighted bipartite graph  $G$ . We have an edge  $e$  between a product  $b \in B$  and a user  $u \in U$  with  $e = (u, b)$  and  $w(e) \in \mathbb{R}_{\geq 0}^k$  where  $k$  is the number of attributes users rate each product on. We will represent the data using an adjacency tensor  $\mathcal{A} \in \mathbb{R}^{|B| \times |U| \times |B| \times |U| \times k}$ . If we look at the structure a frontal slice of  $\mathcal{A}$  we see that  $\mathcal{A}$  has the following block diagonal structure

$$\begin{bmatrix} A & 0 \\ 0 & A^\top \end{bmatrix}.$$

Thus we only have to store  $A \in \mathbb{R}^{|B| \times |U| \times k}$  and achieve a speed up by a factor of 4. Thus, we can define the adjacency tensor of our dataset as

$$\mathcal{A}(u, b, k) = \begin{cases} w((u, b))_k & \text{if } (u, b) \in G \\ 0 & \text{if } (u, b) \notin G \end{cases} \quad (1)$$

Alternating Least Squares (ALS), the traditional method of computing the rank  $R$  CP-Decomposition of a order  $m$  tensor  $\mathcal{X}$ , relies on updating each factor of the matrices  $A_{(n)}$  by solving the following optimization problem [4].

$$A_{(n)}^{(r)} = \underset{A}{\operatorname{argmin}} \left\| \mathcal{X}_{(n)} - H_{(n)} A^\top \right\|_F^2 \quad (2)$$

where  $\mathcal{X}_{(n)}$  is the mode  $N$  unfolding of the tensor  $\mathcal{X}$  and  $H_{(n)} = A_{(1)} \circ A_{(2)} \circ \dots \circ A_{(n-1)} \circ A_{(n)} \circ \dots \circ A_{(N)}$ . We note that computing  $H_{(n)} A_{(x)}^\top$  is computing our rank  $F$  estimate of  $\mathcal{X}_{(n)}$ . The ALS optimization problem given above has a closed form solution given by

$$\left( \left( H_{(n)}^\top H_{(n)} \right)^{-1} H_{(n)}^\top \mathcal{X}_{(n)} \right). \quad (3)$$

Computing  $\left( H_{(n)}^\top H_{(n)} \right)^{-1}$  is relatively simple but as  $\mathcal{X}$  gets larger computing the decomposition becomes infeasible. If  $\mathcal{X} \in \mathbb{R}^{I_1 \times I_2 \times \dots \times I_m}$  then computing the matricized tensor times Khatri-Rao product (MTTKRP)  $H_{(n)}^\top \mathcal{X}_{(n)}$  requires  $O\left(\prod_{n=1}^N I_n F\right)$  operations.

### III. ALGORITHMS

#### A. Stochastic Gradient Descent

In [3] an algorithm for computing the decomposition by randomly sampling fibers of the tensor  $\mathcal{X}$  is given by the following algorithm.

First, we sample a set of mode  $n$  fibers of  $\mathcal{X}$  where each fibers is a row of the mode- $n$  unfolding of the tensor  $\mathcal{X}$ . We denote the set of mode  $n$  fibers sampled from  $\mathcal{X}$  as  $\mathcal{F}_n \subset \{1, 2, 3, \dots, J_n\}$ . If  $|\mathcal{F}_n| > F$  then we can solve the sketched system of equations,

$$A_{(n)}^{(i)} = \underset{A}{\operatorname{argmin}} \left\| \mathcal{X}_{(n)}(\mathcal{F}_n, :) - H_{(n)}(\mathcal{F}_n, :) A^\top \right\|_F^2 \quad (4)$$

by computing,

$$A_{(n)}^{\top(i)} = H_{(n)}(\mathcal{F}_n, :)^\dagger \mathcal{X}_{(n)}(\mathcal{F}_n, :). \quad (5)$$

We could utilize this method but it does not allow for constraining the factors  $A_{(n)}$ . Instead, we consider the stochastic gradient for each factor matrix given by,

$$G_{(n)}^{(r)} = \frac{1}{|\mathcal{F}_n|} \left( A_{(n)}^{(r)} H_{(n)}^T(\mathcal{F}_n) H_{(n)}(\mathcal{F}_n) - \mathcal{X}_{(n)}(\mathcal{F}_n) H_{(n)}(\mathcal{F}_n) \right) \quad (6)$$

$$G_{(n')}^{(r)} = 0, n' \neq n. \quad (7)$$

We can then adaptively select the step size using adagrad [7].

#### B. Constrained Case

If we want to constrain each factor matrix  $A_{(n)}$  to some convex set  $C$  we can apply proximal gradient descent to the convex  $\infty - 0$  indicator function  $\iota_C$  defined as

$$\iota_C(x) = \begin{cases} 0 & \text{if } x \in C \\ \infty & \text{if } x \notin C \end{cases}. \quad (8)$$

Then on each iteration project  $A_n$  into  $C$  by computing

$$\operatorname{prox}(A_n) = \underset{A}{\operatorname{argmin}} \left( \iota_C(a) + \frac{1}{2} \|A - A_{(n)}\|_2^2 \right) \quad (9)$$

$$= \underset{A \in C}{\operatorname{argmin}} \|A - A_{(n)}\|_F^2. \quad (10)$$

We then set update  $A_n^{(r+1)}$  as

$$A_{(n)}^{(r+1)} \leftarrow \operatorname{prox} \left( A_{(n)}^{(r)} - \alpha G_{(n)}^{(r)} \right) \quad (11)$$

where  $\alpha$  is the step size selected by adagrad.

#### C. Masked-CP Decomposition

If we try to solve

$$\min \left\| \mathcal{X}_{(n)} - H_{(n)} A_{(n)}^\top \right\|_F^2 \quad (12)$$

then the problem is dominated by entries with value 0 in  $\mathcal{X}$  where the review is missing. If we naively solve this problem, we will not get good generalization. Thus we introduce the concept of a mask tensor

$\mathcal{M}$  which has the same dimensions as  $\mathcal{X}$ .  $M[u, v, :] = \mathbf{0}$  if we are missing the review user  $u$  gave product  $v$  and 1 otherwise. We then try to solve the related problem

$$\min \left\| M \circ \left( \mathcal{X}_{(n)} - H_{(n)} A_{(n)}^\top \right) \right\|_F^2 \quad (13)$$

$$= \min \left\| \mathcal{X}_{(n)} - M \circ \left( H_{(n)} A_{(n)}^\top \right) \right\|_F^2. \quad (14)$$

We can then obtain an analogous gradient as before

$$G_{(n)}^{(r)} = \frac{1}{|\mathcal{F}_n|} \left( \left( M^\top(\mathcal{F}_n) \circ \left( A_{(n)}^{(r)} H_{(n)}^\top(\mathcal{F}_n) \right) \right) H_{(n)}(\mathcal{F}_n) - \mathcal{X}_{(n)}(\mathcal{F}_n) H_{(n)}(\mathcal{F}_n) \right) \quad (15)$$

apply the same two algorithm as before. The new optimization problem will no longer penalize entries that we do not have data for so the value of  $\mathcal{X}(u, v, :)$  will no longer influence  $A_{(n)}$  if  $(u, v)$  is not in the dataset.

## IV. METHODOLOGY

### A. Procuring Data

The dataset used for this project is a collection of about 1.5 million reviews of beers over a period of 10 years from the website BeerAdvocate. The dataset is provided by the Stanford SNAP Lab and is hosted on data.world.

The dataset is stored as a CSV file. Each row of the file contains the following columns

- 1) *brewery\_id*: unique numerical identifier for brewery that manufactured reviewed beer
- 2) *brewery\_name*: name of brewery that manufactured reviewed beer
- 3) *review\_time*: time of review (as Unix time)
- 4) *review\_overall*: overall rating of beer by reviewer, on scale of 1 to 5 increments of 0.5
- 5) *review\_aroma*: rating of beer aroma by reviewer, on scale of 1 to 5 increments of 0.5
- 6) *review\_appearance*: rating of beer appearance by reviewer, on scale of 1 to 5 increments of 0.5
- 7) *review\_profilename*: Profile name of reviewer
- 8) *beer\_style*: category of beer (one of 104 possible categories, e.g., “Light Lager”)
- 9) *review\_palate*: rating of beer palate by reviewer, on scale of 1 to 5 increments of 0.5
- 10) *review\_taste*: rating of beer taste by reviewer, on scale of 1 to 5 increments of 0.5
- 11) *beer\_name*: name of reviewed beer
- 12) *beer\_abv*: percent alcohol by volume of beer
- 13) *beer\_beerid*: unique numerical identifier for reviewed beer

### B. Preprocessing Data

As noted in II.5, the dataset used for the problem should describe a bipartite graph  $G$ . For this problem, we state the graph is comprised of disjoint sets  $B$  and  $U$ , which are the beers and users, respectively, whose information is stored in the dataset.

1) *Summary Queries*: To interpret the raw data into a bipartite graph, the data is organized into several groups using SQL’s built-in GROUP BY queries. The following queries are written sequentially,

- Group the data by identifier of the reviewed beer (*beer\_beerid*) and store the number of distinct reviewers who have reviewed that beer. This gives the set of beers  $B$ .
- Group the data by the profile name of the reviewer (*reviewer\_profilename*) and store the number of distinct beers they have reviewed. This gives the set of users  $U$ .
- For each beer  $b$  and reviewer  $u$ , if  $u$  has reviewed  $b$ , return  $u$ ’s most recent of review of  $b$ . Return all five attributes reviewed. This gives the set of all edges in the bipartite graph  $(b, u)$  where  $b \in B$  and  $u \in U$ , and where the ratings for each attribute comprise the edge weight tuple.

2) *Minimum Degree Requirements*: Due to the computational requirements of handling the entire set  $U$  and  $B$  we restrict the sets to product nodes with degree  $\geq 30$  and users with degree  $\geq 20$ . We ended up with a total of 1,235,925 reviews (78%), 7,694 users (23%) and 7,136 products (11%).

3) *Transformation into Tensor*: Once the data is filtered through SQL we employ a 80-10-10 train, test, and validation split on the data. From this we create two tensors  $\mathcal{A}$ , the adjacency tensor consisting of beers and users and  $\mathcal{M}$  the mask tensor which denotes whether a review is present or not.

### C. Method

Given our two tensors  $\mathcal{A}$  and  $\mathcal{M}$ , we employ masked CP Decomposition on these two tensors. We test three different models with various ranks. We use the following general algorithm:

---

**Algorithm 1:** Computing the Estimate  $E(u, b)$

---

```

1: procedure DECOMP( $\mathcal{A}, E', R$ )  $\triangleright$  Given the adjacency tensor  $\mathcal{A}$ , validation set  $E'$  and max rank  $R$ 
2:    $r \leftarrow 1$ 
3:   while  $r \leq R$  do
4:     Initialize feature matrices  $A_{(n)}$  randomly
5:      $\mathcal{T} \leftarrow \text{Decomp}(\mathcal{A}, \mathcal{M}, r, \mathbf{A})$ 
6:     Compute  $\frac{1}{|E'|} \sum_{(u,b) \in E'} \|\mathcal{A}_{ub} - \mathcal{T}_{ub;r}\|^2$ 
7:      $R \leftarrow r$  if error is less than previous error.
8:   end procedure

```

---

Our ranks are in the range  $[1, 25]$  and Decomp function uses our Masked-CP decomposition. We use two versions of Decomp: MCP-Decomposition with SGD and MCP-Decompistion with PGD. We run until the relative iterative error is  $\leq 10^{-3}$  or up to 500 iterations and sampled 12500 fibers. To achieve faster results, we use multiprocessing on our three separate models for each rank.

### D. System Information

All experiments were performed on High-Ram Google Colab instances using python3. The tensor decomposition was written for this project but with basic tensor code using tensorly [6]. We also utilized numpy and the multiprocessing library to compute 3 decomposition at a time. We found any more would be constrained by the colab environment's 58 GB memory limitation.

### E. Comparison Model

Here we introduce a simple comparison model that we will compare our tensor factorization method to. Let  $U(u) \in \mathbb{R}^k$  be the average review user  $u$  gave in the training set and let  $B(b)$  be the average review product  $b$  has in the training set. Then we let our estimate  $\tilde{E}(u, b) = \lambda U(u) + (1 - \lambda)B(b)$  where  $\lambda$  is chosen from the validation set. It is clear that the comparison model is a tensor of rank at most  $2k$ .

## V. RESULTS AND DISCUSSION

In this section, we present our results from our tensor methods. We measure the error as a function of the rank of the decomposition. We compute the in-sample error  $E_{in}$  as well as the validation error  $E_{val}$  and take an average amongst our three models. This is shown in Figure 1 and Figure 2. In the figures, there is a similar trend in using both PGD and SGD. However, we find that PGD has a slight underhand producing an  $E_{test} = 0.0906$ . The best rank decomposition is of Rank 4. For rank  $> 4$ , the error starts to increase.

For our comparison model, we compute  $\lambda$  based on our validation set thus minimizing the validation error. Applying this value of  $\lambda$ , we get the following  $E_{test} = 0.0534, E_{val} = 0.0615, E_{in} = 0.0613$ . The comparison model here outperforms our tensor model by a factor  $\approx 1.8$ .




Fig. 1. Masked CP-Decomposition with Proximal Gradient Descent




Fig. 2. Masked CP-Decomposition with Stochastic Gradient Descent

### A. Discussion

Although there was success in modeling the problem, the in-sample error,  $E_{in}$ , is still relatively high. Due to limitations of computational resources, this was the best our model could perform. For our sparse data set with a large number of nodes, the tensor scales as a function of  $O(MNk)$  and the cost per iteration of gradient descent is  $O(F|\mathcal{F}_n|I_n)$ . Other randomized CP-Decomposition algorithms could perform better such as second order methods such as those introduced in [5]. Additionally, if more computational resources were feasible, setting our threshold lower/running the algorithm for longer may provide for a smaller  $E_{in}$  yielding a better result.

## VI. CONCLUSIONS

We have introduced a new model for predicting vector weighted edges in bipartite graphs based on computing the latent factors through CP decomposition. Our method is computationally expensive and suffers from a high in sample error when compared to a simple comparison model of rank  $\leq 10$ . It is still possible with more computational resources, time or other algorithms that this method could out substantially out perform the comparison model.

Future work could comprise of non-uniform sampling methods or utilizing a non-variable sampling rate  $|\mathcal{F}_n|$ , for instance we could sample fibers corresponding to each node proportionally the degree of the node. We could also utilize the sparsity of both  $M$  and  $\mathcal{X}$  to achieve a speedup in computing the decomposition. Investigating other decomposition such as one of the variants of the Tucker decomposition [2]. Exploring more ways to constrain each  $A_{(n)}$  may also yield good generalization properties if  $E_{in}$  is made sufficiently small. Other dataset with not just vector edges more general tensor would also work with this frame work, though they would require an tensor of order 4 or higher and would be constrained by the computational difficulty.

## REFERENCES

- [1] Y. Koren, R. Bell and C. Volinsky, "Matrix Factorization Techniques for Recommender Systems," in *Computer*, vol. 42, no. 8, pp. 30-37, Aug. 2009, doi: 10.1109/MC.2009.263.
- [2] Kolda, Tamara G. and Bader, Brett W (2009). Tensor Decompositions and Applications. *SIAM Review*, 51(3), 455-500.
- [3] Frolov, Evgeny, Huang, Ibrahim, Wai 2019]p2 Xiao Fu, Cheng Gao, Kejun Huang, Shahana Ibrahim, Hoi-To Wai, 2019 Block-Randomized Stochastic Proximal Gradient for Low-Rank Tensor Factorization.
- [4] Yipeng Liu, Jiani Liu, Zhen Long, Ce Zhu, Yipeng Liu, Jiani Liu, Zhen Long, Ce Zhu, Tensor Decomposition, Tensor Computation for Data Analysis, 10.1007/978-3-030-74386-4, (19-57), (2022).
- [5] Gittens, A., Aggour, K., & Yener, B. (2020). Adaptive Sketching for Fast and Convergent Canonical Polyadic Decomposition. In *Proceedings of the 37th International Conference on Machine Learning* (pp. 3566–3575). PMLR.
- [6] Jean Kossaifi, Yannis Panagakis, Anima Anandkumar and Maja Pantic, TensorLy: Tensor Learning in Python, *Journal of Machine Learning Research*, Year: 2019, Volume: 20, Issue: 26, Pages: 16. <http://jmlr.org/papers/v20/18-277.html>.
- [7] John Duchi, Elad Hazan, Yoram Singer (2011). Adaptive Subgradient Methods for Online Learning and Stochastic Optimization. *Journal of Machine Learning Research*, 12(61), 2121-2159.